

GANs for Data Augmentation

Eric Gentry and Jiahua You

[Link to report](#);
[link to github repo](#)

Gentry and You -- CMPS 292C -- Winter 2018

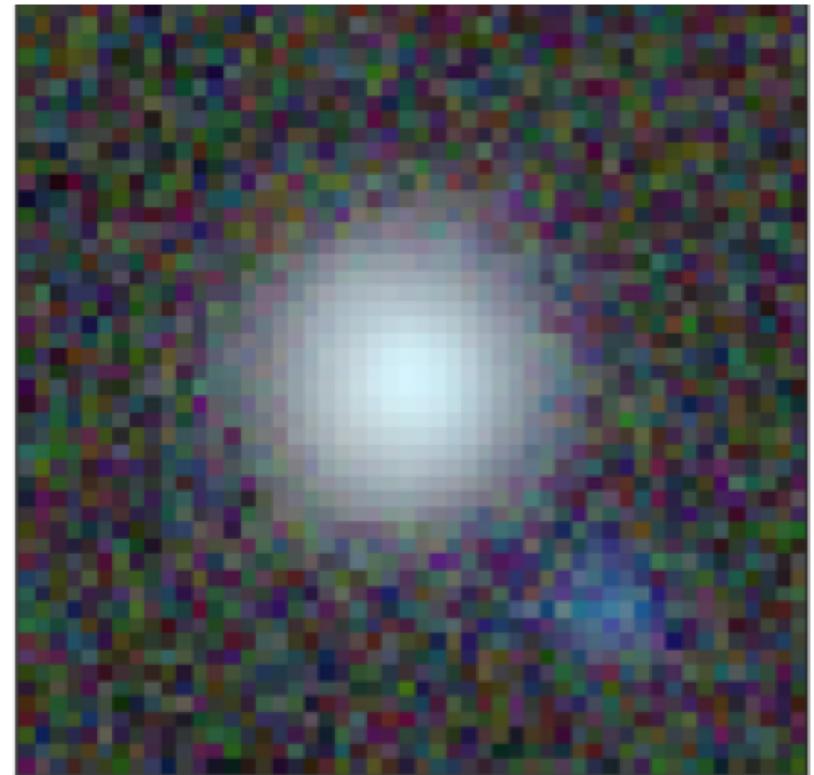
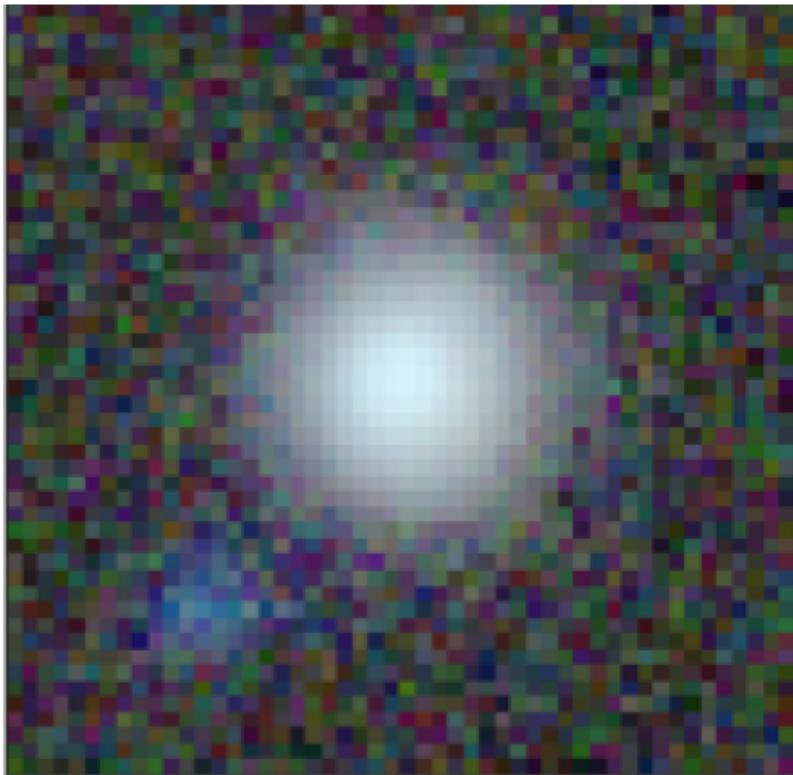
Goal:

1. Split the data into **training and testing sets**.
2. Build one **classifier that is trained on the real data**, using traditional data augmentation transformations (reflections, translations, etc). This will be the **benchmark**.
3. **Build a conditional GAN**, which accepts the metadata features “distance of galaxy” and “mass of galaxy”, and generates “realistic” images for that distance and mass. Train the GAN and then freeze its weights.
4. **Use this GAN to create training images for another classifier**, *using the same architecture* as in Step 2. This classifier is trained on purely generated images, and validated on purely real images.
5. **Compare results** from classifier trained using traditional data augmentation (Part 2) and the classifier trained using GAN augmentation (Part 4).

Reminder: What is data augmentation

Transformations that don't change the true label

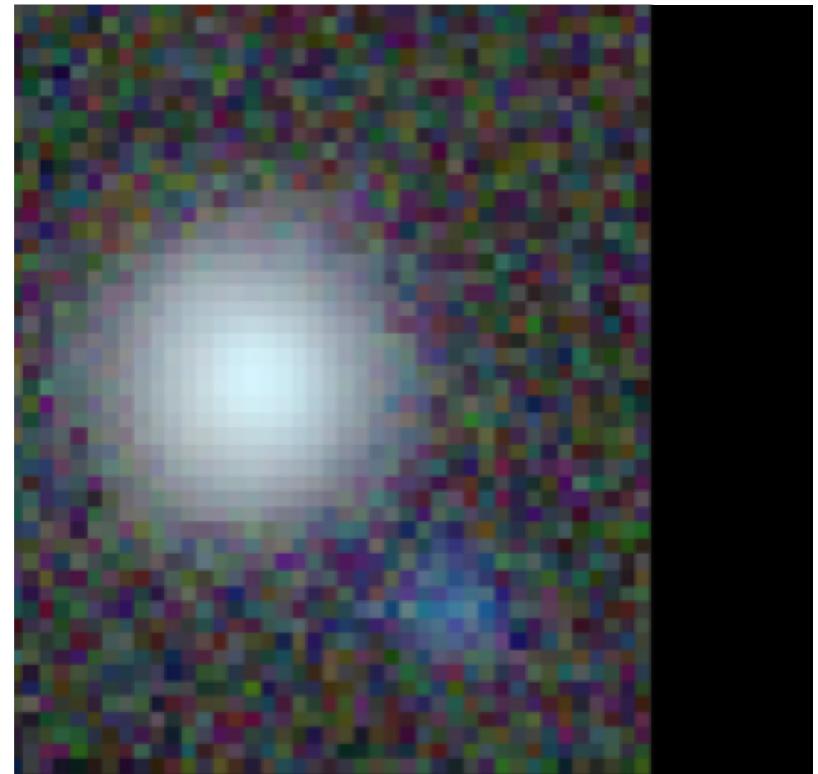
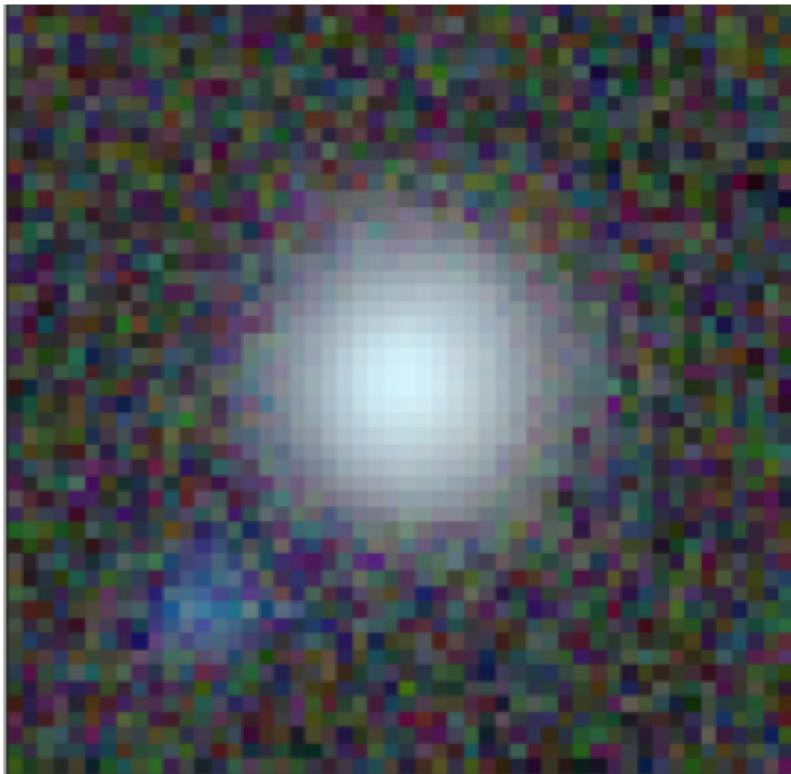
Example: Horizontal reflection



Reminder: What is data augmentation

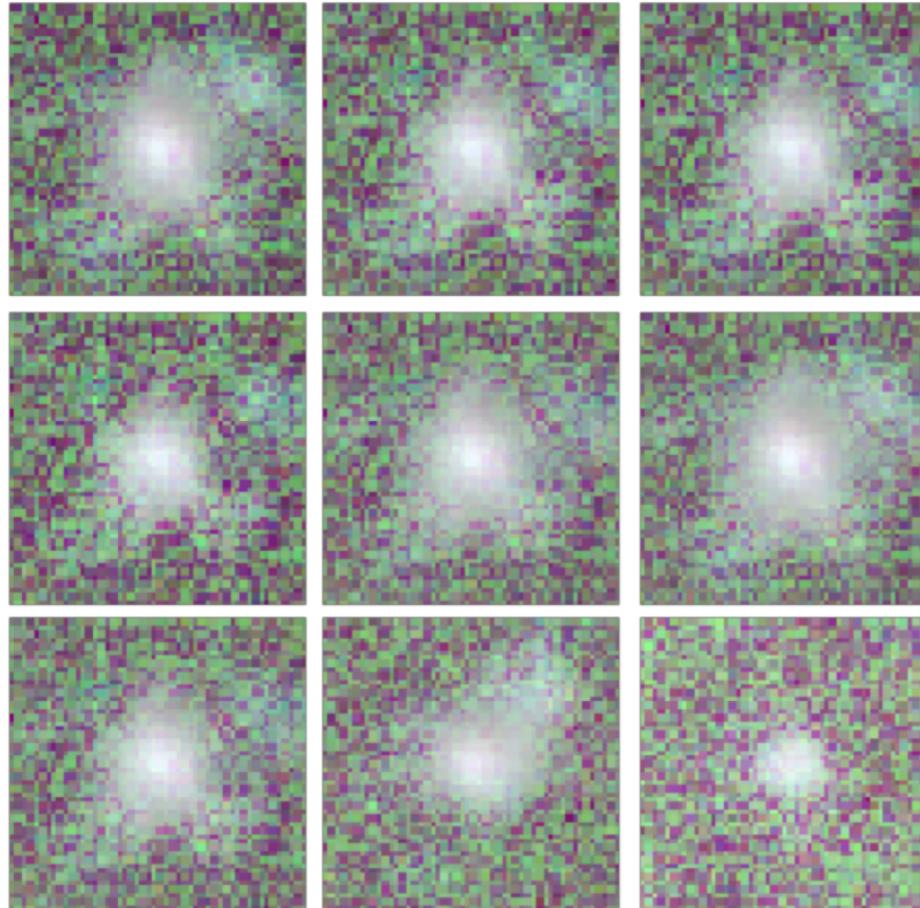
Transformations that don't change the true label

Example: Horizontal *translation* (exaggerated)



Would be much nicer if a GAN could just generate images for us

But a **reminder**:
We want useful
pictures, not just
pretty pictures



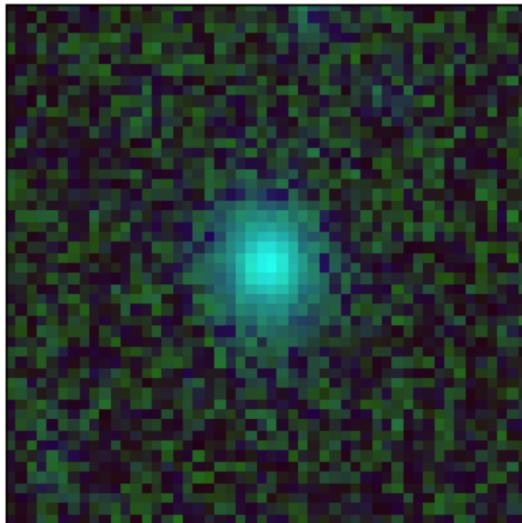
Task: classify galaxy images

Data: 2000 images * 3 channels * (50 px)²

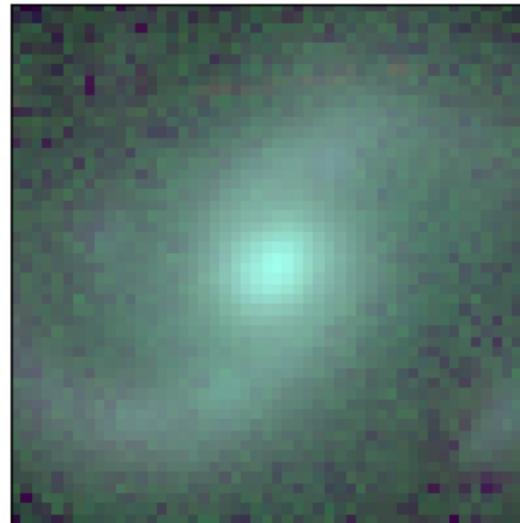
Labels:

- Distance, Mass
- “nearby low-mass galaxy” (boolean; class balance= 2:1 (False:True))

Label = True

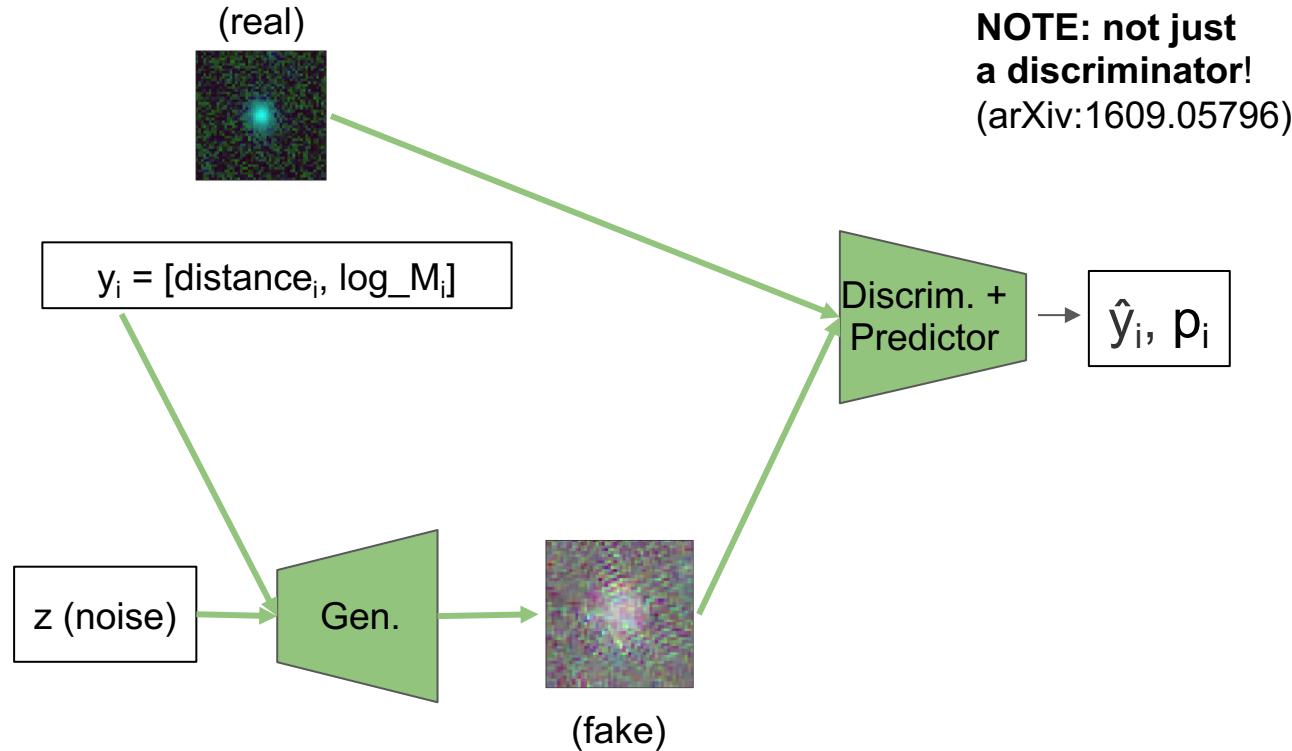


Label = False

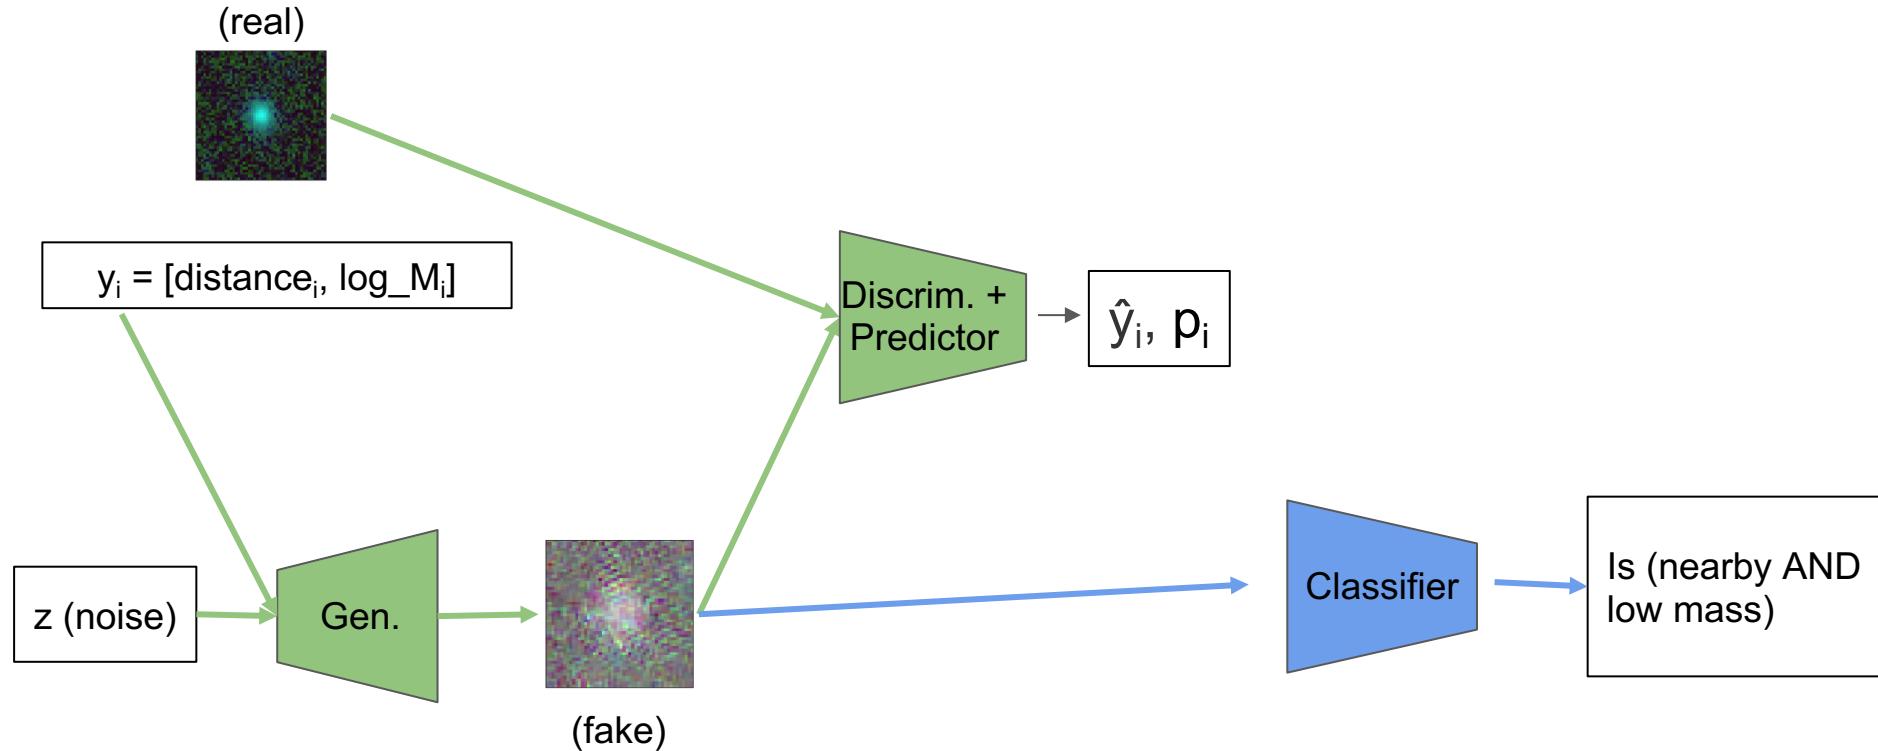


Same distance,
But 30x more massive

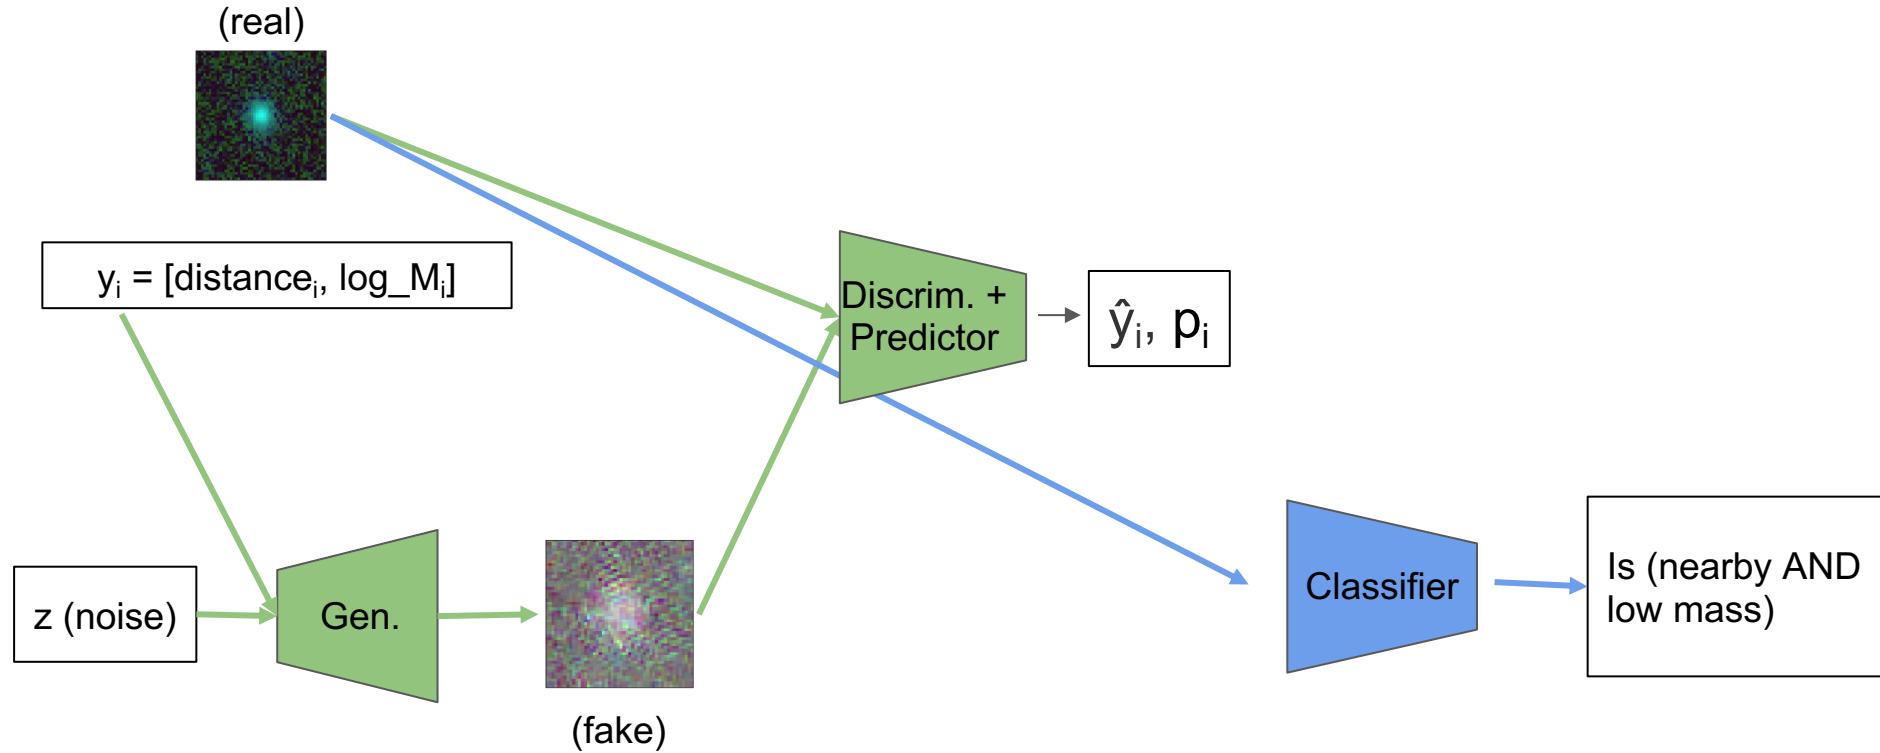
Architecture: Overview (training GAN)



Architecture: Overview (GAN *training* Classifier)

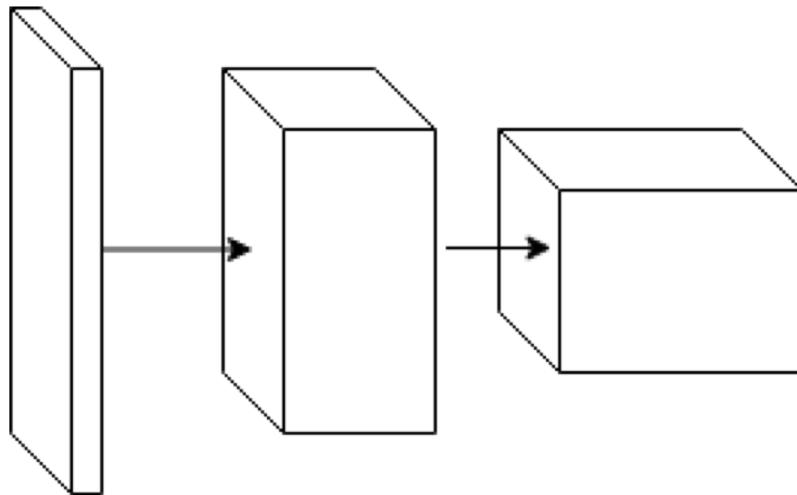


Architecture: Overview (Classifier *validation*)



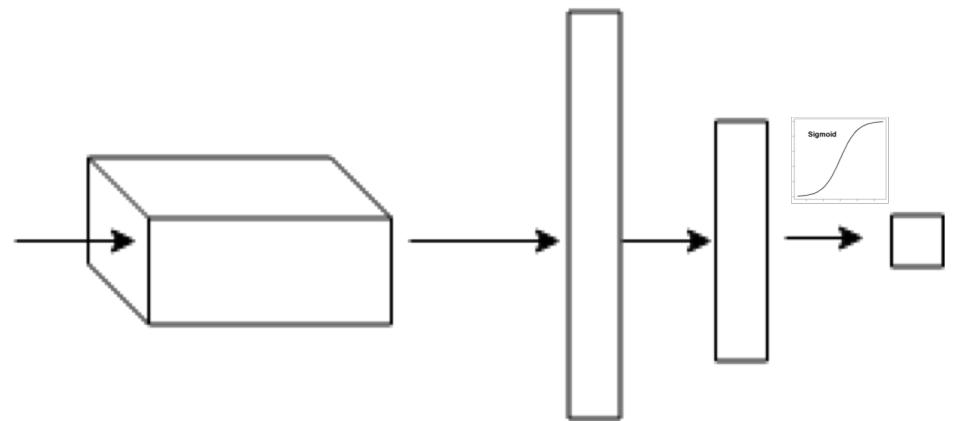
Architecture: Classifier

Conv Layers



RELU activation,
MaxPool,
Dropout between layers

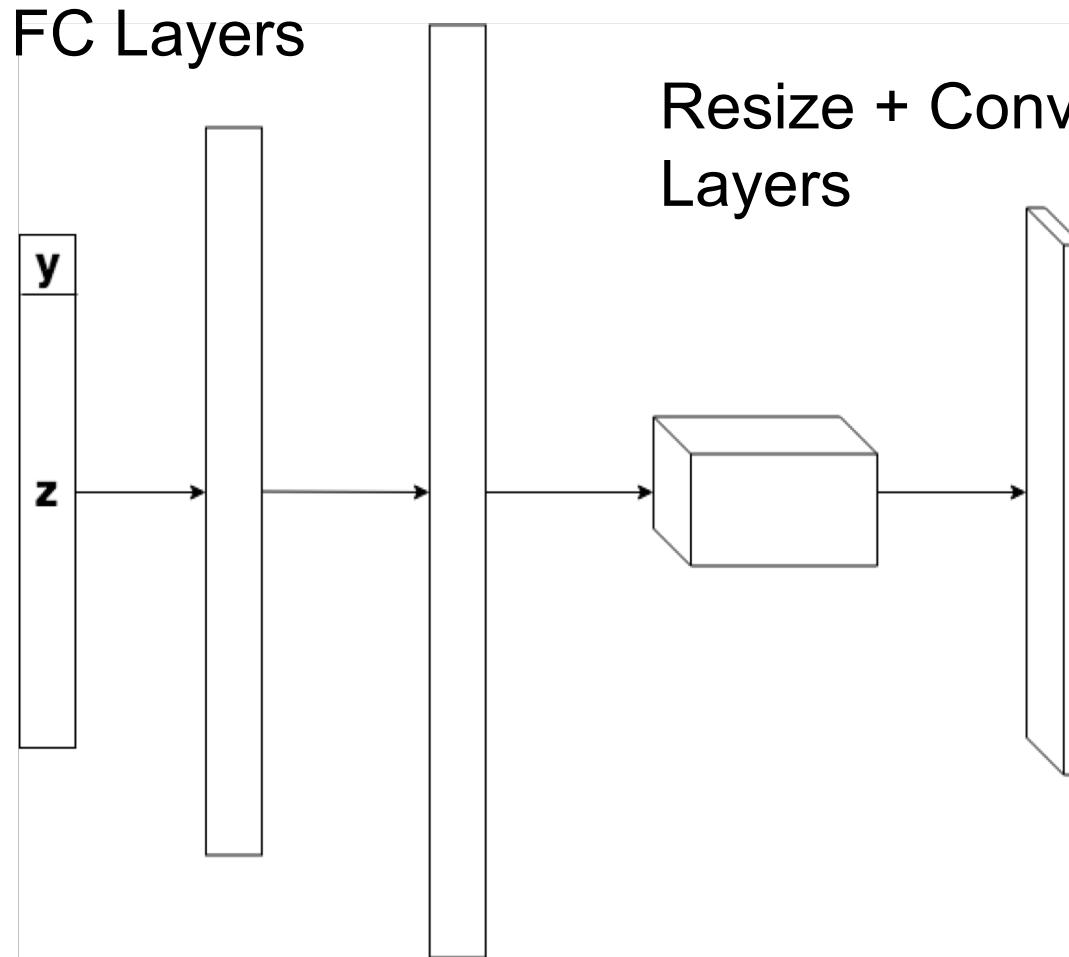
FC Layers



RELU activation,
Except sigmoid for last layer

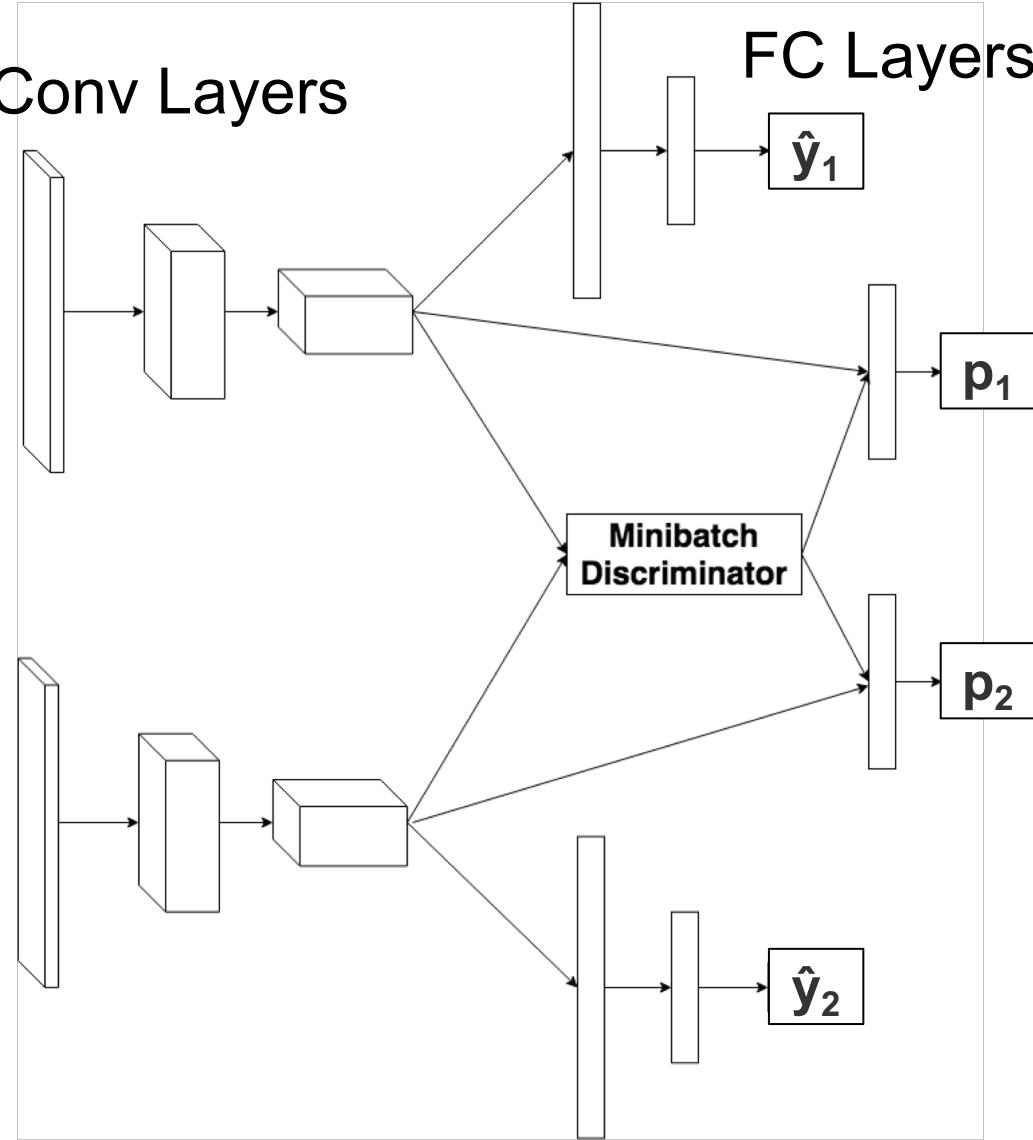
Loss function: binary cross-entropy

Architecture: GAN --- Generator

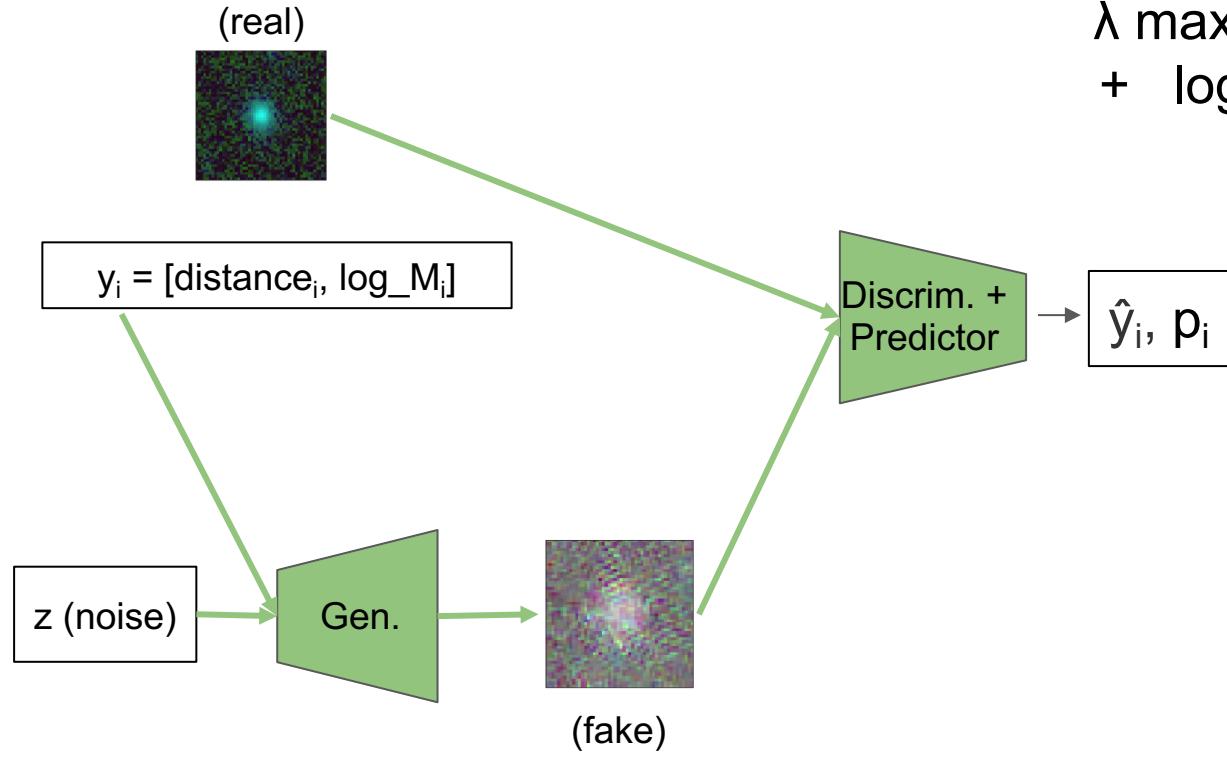


Note: not deconv. layers!
(due to checkerboard artifacts)

Architecture: GAN --- Discriminator/Predictor



Architecture: GAN --- Overall



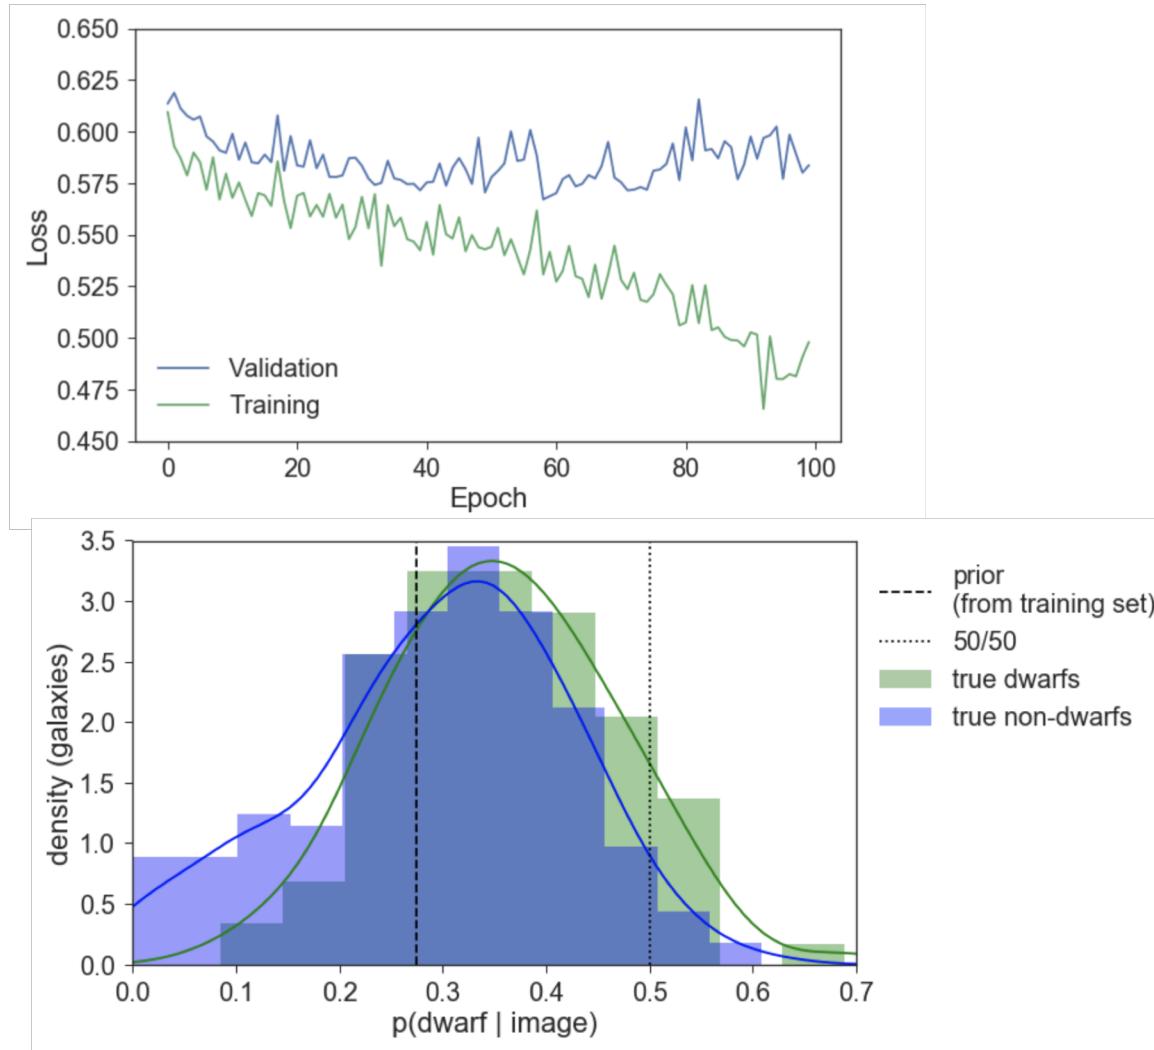
Discrim./Predictor Loss:

$$\lambda \max[0, \ell^2(\hat{y}_{\text{real}}) - \ell^2(\hat{y}_{\text{fake}})] + \log(p_{\text{real}}) + \log(1 - p_{\text{fake}})$$

Generator Loss:

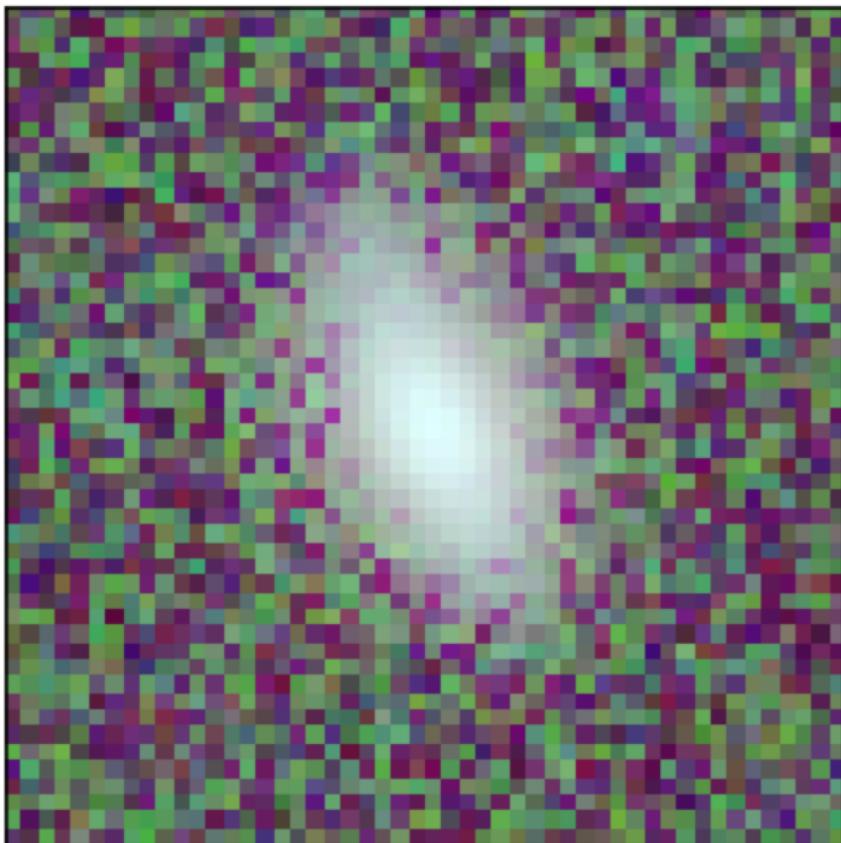
$$\lambda \ell^2(\hat{y}_{\text{fake}}) + \log(p_{\text{fake}})$$

Results - Classifier Baseline

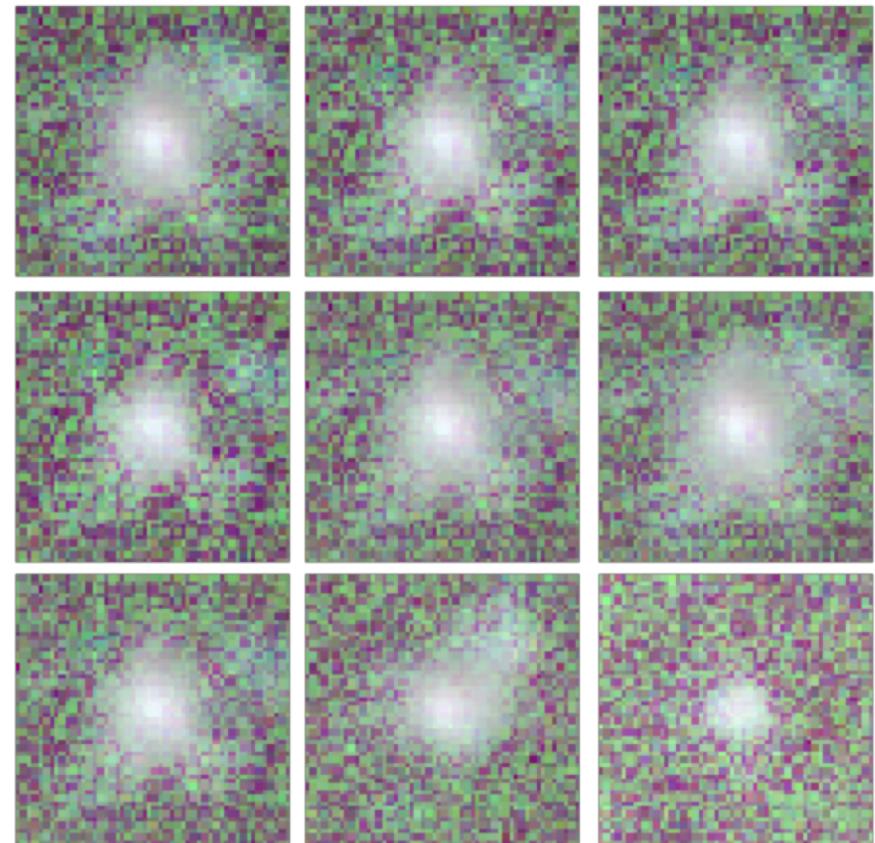


Results: GAN generated images are decent

Real



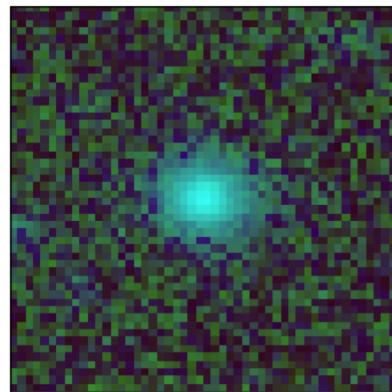
GAN-generated



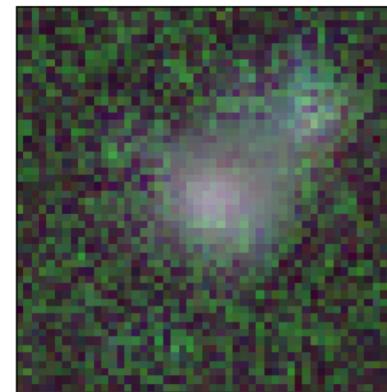
Results: GAN doesn't always capture conditional label

log_Mass=8.5

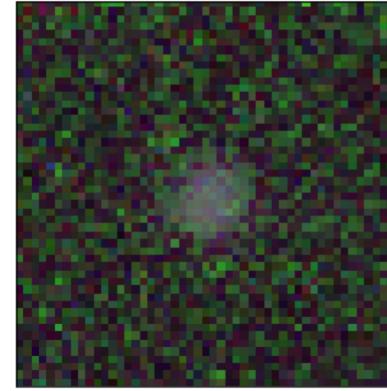
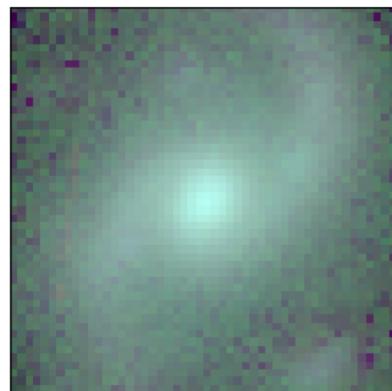
Real



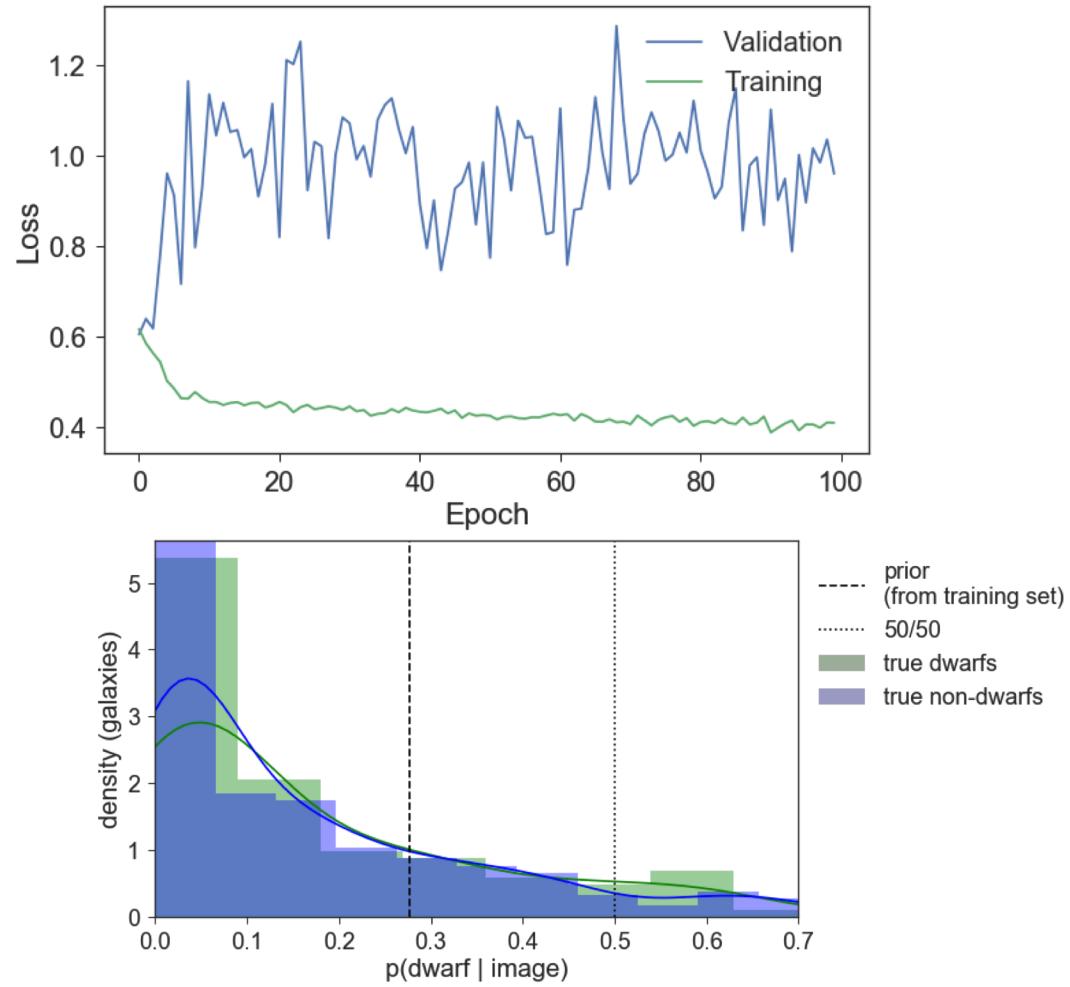
GAN-generated



log_Mass=10.1

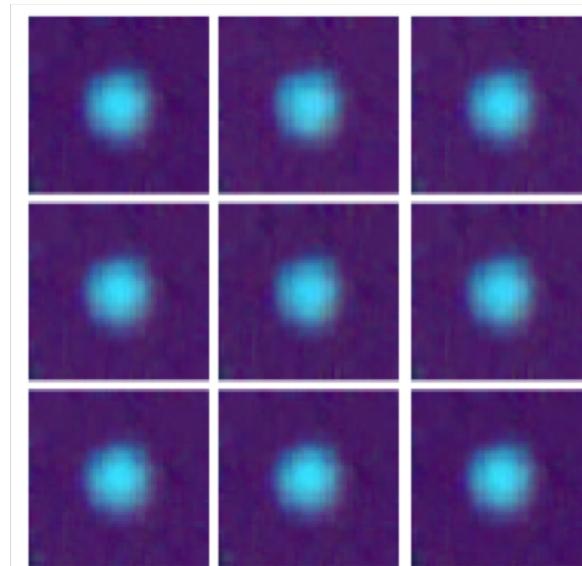


Results - Training Classifier on GAN images *hurts* performance



Problems encountered

1. Difficult to generate good images given **continuous conditional values**
 - a. Added *predictor* in addition to *discriminator* (see [Ravanbakhsh et al. 2016](#))
2. Generator **mode collapse**
 - a. Partial solution: use *minibatch discrimination*, which adds a loss for correlated generated samples (see Salimans et al. 2016)
 - b. Requires balancing an L2 loss and a cross-entropy loss



9 “independent” samples from
an *overtrained GAN*
(without minibatch discrim.)

Conclusions

Creating Conditional-GAN is non-trivial, even for “simple” images

- Needed to hand-tune: architecture, loss functions, optimizer parameters

Training on C-GAN images decreased performance

- Probably because it doesn't capture label-dependence properly
- If label dependence were fixed, maybe we'd get better performance by first training on GAN images, and then training on real images